

PREFERENCE LEARNING PHYSICAL SIDE EFFECTS: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study human preference learning for robots. The motivating bet is: Learn preferences over physical side effects that are absent from task success labels.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For human preference learning for robots, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2024a;b;c; 2018; 2006; 2017). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.532	0.468
observed_only_wam	hidden_mode_shift	0.160	0.840
augmentation_as_data_wam	nominal	0.575	0.425
augmentation_as_data_wam	hidden_mode_shift	0.258	0.742
scalar_uncertainty_wam	nominal	0.657	0.343
scalar_uncertainty_wam	hidden_mode_shift	0.312	0.688
proposed_branch_mechanism	nominal	0.797	0.203
proposed_branch_mechanism	hidden_mode_shift	0.535	0.465

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Preference Learning Physical Side Effects argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Are physically embodied social agents better than disembodied social agents?: The effects of physical embodiment, tactile interaction, and people’s loneliness in human-robot interaction, 2006. <https://doi.org/10.1016/j.ijhcs.2006.05.002>.
- A user study on personalized stiffness control and task specificity in physical human-robot interaction, 2017. <https://doi.org/10.3389/frobt.2017.00058>.
- Preference-based assistance prediction for human-robot collaboration tasks, 2018. <https://doi.org/10.1109/iros.2018.8593716>.
- Apricot: Active preference learning and constraint-aware task planning with llms, 2024a. <http://arxiv.org/abs/2410.19656v1>.
- Unified learning from demonstrations, corrections, and preferences during physical human-robot interaction, 2024b. <https://doi.org/10.1145/3623384>.
- Drivellava: Human-level behavior decisions via vision language model, 2024c. <https://doi.org/10.3390/s24134113>.